T Q1 2025: Foundations & PyTorch Mastery

Goal: Build strong fundamentals in ML, PyTorch, and Deep Learning.

ML & Python for Al

- Revise Python (NumPy, Pandas, Matplotlib, Scikit-Learn)
- V Study core ML concepts (Supervised, Unsupervised Learning)
- V Implement Regression, Classification models in Scikit-Learn
- Learn probability, statistics, and linear algebra for ML
 Courses:
- Python for Data Science (Udemy)
- Machine Learning Specialization (Andrew Ng)

PyTorch & Deep Learning Fundamentals

- Master PyTorch basics (Tensors, Autograd, Neural Networks)
- V Implement Feedforward & Convolutional Neural Networks (CNNs)
- Train models on MNIST & CIFAR-10 datasets
- V Understand loss functions, optimizers (SGD, Adam)
 - **⊗** Courses:
- Deep Learning with PyTorch (Udacity)
- Fast.ai Deep Learning Course
- Udemy (Daniel Bourke) Pytorch

Advanced Neural Networks & Deployment Basics

- V Learn RNNs, LSTMs, Transformers
- V Study activation functions, batch normalization, dropout
- Introduction to model deployment (TorchServe, Flask)
- Mini Project: Build & deploy an image classifier
 Courses:
- Deep Learning Specialization (Coursera)

📆 Q2 2025: Specializing in NLP & Computer Vision

Goal: Dive deep into NLP and Vision.

NLP Basics & Transformers

- V Learn Text Preprocessing (Tokenization, Embeddings)
- Study RNNs, LSTMs for NLP tasks
- Work with Hugging Face Transformers library
- V Fine-tune BERT on sentiment analysis
 - **⊗** Courses:
- NLP Specialization (Coursera)
- Hugging Face's NLP Course

Computer Vision (CNNs, Object Detection)

- Train CNN models for image classification
- V Study object detection (YOLO, Faster R-CNN)
- **V** Experiment with segmentation models (U-Net)
- Mini Project: Implement real-time object detection
 - **♦** Courses:
- Convolutional Neural Networks (Coursera)
- PyTorch Vision Tutorials

Large Language Models & Fine-Tuning

- V Study Transformer architectures (GPT, T5)
- V Fine-tune LLaMA, GPT-4 on custom datasets
- V Learn Retrieval-Augmented Generation (RAG)
- V Deploy a chatbot using a fine-tuned LLM
 - **Ourses:**
- Hugging Face's LLM Course
- Generative AI with LLMs (DeepLearning.AI)

📆 Q3 2025: Multimodal Al & Advanced Al Research

Goal: Work on Vision + NLP (Multimodal AI) and advanced AI models.

Multimodal Learning (Vision + NLP)

- ✓ Study CLIP, DALL·E, Flamingo architectures
- - **Resources**:
- Hugging Face's Multimodal Al Course

Generative AI & Diffusion Models

- V Study Stable Diffusion & MidJourney
- Z Experiment with DreamBooth & Text-to-3D
- Mini Project: Fine-tune a text-to-image model
- Stable Diffusion Paper
- Hands-On Generative AI with Stable Diffusion

MLOps & Al Model Deployment

- V Learn MLOps fundamentals (CI/CD, Model Serving)
- V Study deployment platforms (TensorRT, Triton)
- Mini Project: Deploy an LLM-powered API
 - **⊗** Courses:
- MLOps Specialization (Coursera)

📆 Q4 2025: Research, Competitions & Specialization

Goal: Work on real-world applications & contribute to Al research.

Research & Cutting-Edge Al

- V Follow research papers on NeurlPS, ICML, CVPR
- Implement papers from Papers with Code
- Contribute to open-source Al projects
- Two Minute Papers YouTube

Al Hackathons & Kaggle Competitions

- V Join Kaggle competitions in NLP & CV
- Participate in Al hackathons
- V Publish a blog on Medium/Kaggle Notebooks
 - **⊗** Resources:
- Kaggle Competitions

Final Capstone Project

- Work on a final AI project (choose one):
 - Al-powered chatbot (LLM + RAG)
 - Multimodal search engine (CLIP-based)
 - Al-generated artwork (Stable Diffusion)

- Deploy project on GitHub, Hugging Face Spaces, or a personal website
 Resources:
- FastAPI for AI Deployment

Quarter	Focus Area
Q1	ML, PyTorch, Deep Learning Basics
Q2	NLP, Computer Vision, LLM Fine-Tuning
Q3	Multimodal AI, Generative AI, Deployment
Q4	Al Research, Kaggle, Final Capstone Project